

# Theoretical Justification of Deep Learning

Generalization and Optimization

Oualid Missaoui

# Agenda


- Standard/classic statistical learning theory (SLT)
- Double Descent and limitations of classic SLT
- Benign Overfitting in interpolating shallow models
  - Linear Models
  - Kernel Methods
  - Boosting
  - Random Forests
- Theoretical Justifications of linear interpolators
- Linearization of Neural Networks
- Optimization: SGD Implicit Bias
- References

# Classic/Standard Statistical Learning Theory

	Finite $\mathcal{H}$	Infinite $\mathcal{H}$
Realizable	$\mathfrak{R}(g) \leq \frac{1}{N} [\log  \mathcal{H}  + \log \frac{1}{\delta}]$	$\mathfrak{R}(g) < \frac{VC(\mathcal{H}) \log 2N + \log \frac{1}{\delta}}{N}$
Agnostic	$\mathfrak{R}(g) < \hat{\mathfrak{R}}(g) + \sqrt{\frac{1}{2N} \log \frac{2 \mathcal{H} }{\delta}}$	$\mathfrak{R}(g) < \hat{\mathfrak{R}}(g) + \sqrt{\frac{VC(\mathcal{H})(\log \frac{2N}{VC(\mathcal{H})}) + \log \frac{4}{\delta}}{N}}$

where  $VC(\mathcal{H})$  is the Vapnik-Chervonenkis dimension of the hypothesis space  $\mathcal{H}$ . What is  $VC(\mathcal{H})$  dimension ?

- A dichotomy of a set  $S$  is a partition of  $S$  into two disjoint subsets
- A set of instances  $S$  is shattered by a hypothesis space  $\mathcal{H}$  if and only if for every dichotomy of  $S$  there exists some hypothesis in  $\mathcal{H}$  consistent with this dichotomy.
- The Vapnik-Chervonenkis dimension  $VC(\mathcal{H})$  of hypothesis space  $\mathcal{H}$  defined over instance space  $\mathcal{X}$  is the size of the largest finite subset of  $\mathcal{X}$  shattered by  $\mathcal{H}$ . If arbitrary large finite sets of  $\mathcal{X}$  can be shattered by  $\mathcal{H}$ , then  $VC(\mathcal{H}) = \infty$




# VC Dimension

Hypothesis class	VC dimension (approx/exact)
Half-line thresholds on $\mathbb{R}$ , $h_a(x) = \mathbf{1}\{x \geq a\}$	1
Intervals on $\mathbb{R}$ (positive on $[a, b]$ )	2
Union of $k$ intervals on $\mathbb{R}$	$2k$
Homogeneous hyperplanes in $\mathbb{R}^d$ (through origin)	$d$
Affine hyperplanes in $\mathbb{R}^d$ (with bias)	$d + 1$
Axis-aligned rectangles in $\mathbb{R}^d$	$2d$ ( $\Rightarrow$ in $\mathbb{R}^2$ : 4)
Euclidean balls in $\mathbb{R}^d$	$d + 1$
Decision stumps on $d$ real features	$d$
Depth- $L$ axis-aligned decision trees	$O(2^L \log 2^L)$ (grows with #leaves)
1-Nearest Neighbor in $\mathbb{R}^d$ (prototypes unconstrained)	$\infty$
Degree- $p$ polynomial separators in $\mathbb{R}^d$	$\binom{d+p}{p}$
ReLU (piecewise-linear) NN with $W$ weights	$\Theta(W \log W)$ (tight up to constants)
Linear separators with margin $\gamma$ on radius- $R$ data	$\lesssim (R/\gamma)^2$ (dimension-free bound)

VC measures worst-case shattering power; higher VC implies a greater sample requirement or stronger regularization for good generalization.

# Double Descent (1)



**Fig. 1.** Curves for training risk (dashed line) and test risk (solid line). (A) The classical U-shaped risk curve arising from the bias-variance trade-off. (B) The double-descent risk curve, which incorporates the U-shaped risk curve (i.e., the “classical” regime) together with the observed behavior from using high-capacity function classes (i.e., the “modern” interpolating regime), separated by the interpolation threshold. The predictors to the right of the interpolation threshold have zero training risk.

# Double Descent (2)

- As model capacity increases from small to moderate, test error first decreases and then increases.
- There is a point where the model exactly fits the training data while test error is near its highest.
- Beyond this point, further increasing capacity causes test error to decrease again.
- Training error keeps decreasing with capacity and remains near zero after exact fit.
- The behavior therefore is not a single U-shape but a “double descent” in test error as capacity grows.
- Models that are far larger than needed to interpolate the data can exhibit lower test error than smaller models.
- The location of the peak and the rate of the second descent depend on the task, data, and model family.

# Double Descent (3)

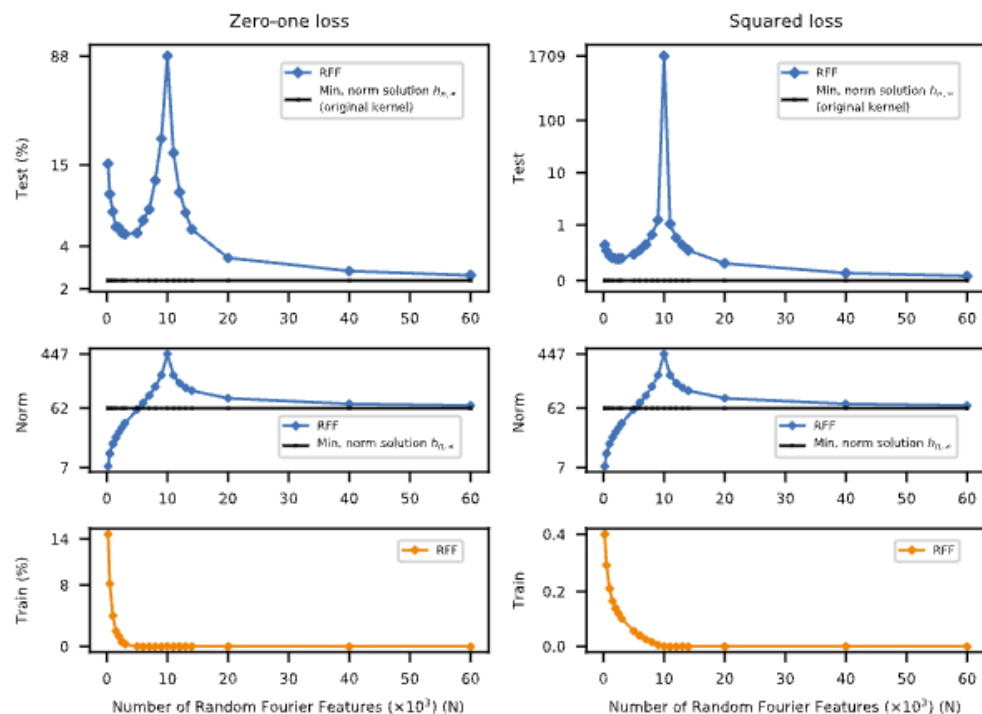


Fig. 2. Double-descent risk curve for the RFF model on MNIST. Shown are test risks (log scale), coefficient  $\ell_2$  norms (log scale), and training risks of the RFF model predictors  $h_{n,N}$  learned on a subset of MNIST ( $n = 10^4$ , 10 classes). The interpolation threshold is achieved at  $N = 10^4$ .




Fig. 3. Double-descent risk curve for a fully connected neural network on MNIST. Shown are training and test risks of a network with a single layer of  $H$  hidden units, learned on a subset of MNIST ( $n = 4 \cdot 10^3$ ,  $d = 784$ ,  $K = 10$  classes). The number of parameters is  $(d + 1) \cdot H + (H + 1) \cdot K$ . The interpolation threshold (black dashed line) is observed at  $n \cdot K$ .




Fig. 4. Double-descent risk curve for random forests on MNIST. The double-descent risk curve is observed for random forests with increasing model complexity trained on a subset of MNIST ( $n = 10^4$ , 10 classes). Its complexity is controlled by the number of trees  $N_{\text{tree}}$  and the maximum number of leaves allowed for each tree  $N_{\text{leaf}}^{\max}$ .

# Double Descent (4)

- The **double-descent pattern** shows up in **all three**: a linear model (with random features), a **random forest**, and a **neural network**.
- For each model family, **training error** falls to (near) zero as complexity increases and stays near zero afterward.
- For each model family, **test error peaks around the interpolation threshold**—the point where the model first exactly fits the training data.
- **Linear model**: increasing the number of random features moves from under- to over-parameterized and produces the same peak-then-drop in test error.
- **Neural network**: growing hidden units/parameters shows the same peak at interpolation followed by a second descent in test error.
- **Random forest**: increasing trees/allowed depth shows the same pattern—test error rises to a peak and then declines with more capacity.
- Across losses (zero-one and squared), the qualitative behavior is the **same** for all three model classes.
- The **peak location** and **descent rate** differ by model family and the chosen complexity knob, but the **overall pattern is consistent**.



# Justification of Double Descent in linear regression (1)

We consider a regression problem where the response  $y$  is equal to a linear function  $\beta = (\beta_1, \dots, \beta_D) \in \mathbb{R}^D$  of  $D$  real-valued variables  $\mathbf{x} = (x_1, \dots, x_D)$  plus noise  $\sigma\epsilon$ :

$$y = \mathbf{x}^* \beta + \sigma\epsilon = \sum_{j=1}^D x_j \beta_j + \sigma\epsilon.$$

Given  $n$  iid copies  $((\mathbf{x}^{(i)}, y^{(i)}))_{i=1}^n$  of  $(\mathbf{x}, y)$ , we fit a linear model to the data only using a subset  $T \subseteq [D] := \{1, \dots, D\}$  of  $p := |T|$  variables.

Let  $\mathbf{X} := [\mathbf{x}^{(1)} | \dots | \mathbf{x}^{(n)}]^*$  be the  $n \times D$  design matrix, and let  $\mathbf{y} := (y^{(1)}, \dots, y^{(n)})$  be the vector of responses. For a subset  $A \subseteq [D]$  and a  $D$ -dimensional vector  $\mathbf{v}$ , we use  $\mathbf{v}_A := (v_j : j \in A)$  to denote its  $|A|$ -dimensional subvector of entries from  $A$ ; we also use  $\mathbf{X}_A := [\mathbf{x}_A^{(1)} | \dots | \mathbf{x}_A^{(n)}]^*$  to denote the  $n \times |A|$  design matrix with variables from  $A$ . For  $A \subseteq [D]$ , we denote its complement by  $A^c := [D] \setminus A$ . Finally,  $\|\cdot\|$  denotes the Euclidean norm.

We fit regression coefficients  $\hat{\beta} = (\hat{\beta}_1, \dots, \hat{\beta}_D)$  with

$$\hat{\beta}_T := \mathbf{X}_T^\dagger \mathbf{y}, \quad \hat{\beta}_{T^c} := \mathbf{0}.$$

Above, the symbol  $^\dagger$  denotes the Moore-Penrose pseudoinverse. In other words, we use the solution to the normal equations  $\mathbf{X}_T^* \mathbf{X}_T \mathbf{v} = \mathbf{X}_T^* \mathbf{y}$  of least norm for  $\hat{\beta}_T$  and force  $\hat{\beta}_{T^c}$  to all-zeros.

**Theorem 1.** Assume the distribution of  $\mathbf{x}$  is the standard normal in  $\mathbb{R}^D$ ,  $\epsilon$  is a standard normal random variable independent of  $\mathbf{x}$ , and  $y = \mathbf{x}^* \beta + \sigma\epsilon$  for some  $\beta \in \mathbb{R}^D$  and  $\sigma > 0$ . Pick any  $p \in \{0, \dots, D\}$  and  $T \subseteq [D]$  of cardinality  $p$ . The risk of  $\hat{\beta}$ , where  $\hat{\beta}_T = \mathbf{X}_T^\dagger \mathbf{y}$  and  $\hat{\beta}_{T^c} = \mathbf{0}$ , is

$$\mathbb{E}[(y - \mathbf{x}^* \hat{\beta})^2] = \begin{cases} (\|\beta_{T^c}\|^2 + \sigma^2) \cdot \left(1 + \frac{p}{n-p-1}\right) & \text{if } p \leq n-2; \\ +\infty & \text{if } n-1 \leq p \leq n+1; \\ \|\beta_T\|^2 \cdot \left(1 - \frac{n}{p}\right) + (\|\beta_{T^c}\|^2 + \sigma^2) \cdot \left(1 + \frac{n}{p-n-1}\right) & \text{if } p \geq n+2. \end{cases}$$

The proof of Theorem 1 is not hard, we give the details in Section 2.2. We now turn to the risk of  $\hat{\beta}$  under a random selection model for  $T$ .

# Justification Double Descent in linear regression (2)




Figure 1: Plot of risk  $\mathbb{E}[(y - \mathbf{x}^* \hat{\beta})^2]$  as a function of  $p$ , under the random selection model of  $T$ . Here,  $\|\beta\|^2 = 1$ ,  $\sigma^2 = 1/25$ ,  $D = 100$ , and  $n = 40$ .

- In this setup with  $n=40$  and  $D=100$ , least-squares regression uses  $p$  features, and as  $p$  increases up to  $n$ , the model overfits and risk worsens—reflecting the classical double descent without the initial descent. Interestingly, when  $p > n$ , the risk decreases again as the model benefits from overparameterization, perfectly fitting the data while becoming less volatile.




Figure 2: Plot of risk  $\mathbb{E}[(y - \mathbf{x}^* \hat{\beta})^2]$  as a function of  $p$ , under the “prescient” selection model of  $T$ . Here,  $\|\beta\|^2 = 1$ ,  $\beta_j^2 \propto 1/j^2$ ,  $\sigma^2 = 1/25$ ,  $D = 100$ , and  $n = 40$ .

- The risk curve shows a “sweet spot” at low feature counts, resembling traditional descent, with higher risk in the interpolation regime. This happens because in a scientific feature selection setting, one cannot hand-pick informative features—so interpolation only helps when feature choice isn’t biased. If features can be cherry-picked, good performance is achievable even with few variables.

# Minimum-norm solution in the under-determined/overparameterized/interpolating regime

**Setting.** Linear regression with more parameters than samples:

- Data matrix  $X \in \mathbb{R}^{n \times d}$  with  $d > n$ ; targets  $y \in \mathbb{R}^n$ .
- Interpolation means there exist  $w \in \mathbb{R}^d$  such that  $Xw = y$  (infinitely many solutions).

**Main result (least-squares).** Among all interpolating solutions,

$$w^\dagger = \arg \min_{w: Xw=y} \|w\|_2 = X^+ y = X^\top (XX^\top)^{-1} y,$$

where  $X^+$  is the Moore–Penrose pseudoinverse.

Equivalently,

$$w^\dagger = \lim_{\lambda \downarrow 0} \arg \min_w \|Xw - y\|_2^2 + \lambda \|w\|_2^2.$$

**How it's selected in practice (implicit bias).**

Gradient Descent on squared loss, initialized at  $w_0 = 0$  (and step size  $< 2/\|X\|_2^2$ ), converges to  $w^\dagger$  among all exact-fit solutions.

**Kernel/RKHS analogue (ridgeless).**

For kernel matrix  $K = X_\phi X_\phi^\top$ :

$$f^\dagger(\cdot) = \arg \min_{f: f(X)=y} \|f\|_{\mathcal{H}} \quad \Rightarrow \quad f^\dagger(x) = k(x, X)^\top K^+ y.$$

**Why it matters.** The minimum-norm (a.k.a. ridgeless) interpolant explains **implicit regularization**, connects to **double-descent/benign overfitting**, and often generalizes when signal aligns with the data span.

**Notation.**  $n$ : samples,  $d$ : features,  $X$ : design matrix,  $y$ : responses,  $w$ : parameters,  $\|\cdot\|_2$ : Euclidean norm,  $X^+$ : pseudoinverse,  $K$ : kernel Gram matrix,  $\mathcal{H}$ : RKHS



## References.

- Hastie, Montanari, Rosset, Tibshirani (2019), *Surprises in High-Dimensional Ridgeless Least Squares*, PNAS.
- Bartlett, Long, Lugosi, Tsigler (2020), *Benign Overfitting in Linear Regression*, PNAS.
- (GD implicit bias) Gunasekar et al. (2018), *Implicit Bias of Gradient Descent on Linear Convolutional Networks*; standard linear case yields the minimum- $\ell_2$  interpolant.

# Bridging to Neural Tangent Kernel (NTK)

- We've seen that **linear models** and **kernel methods** can **interpolate** the data, and gradient descent picks the **minimum-norm** (ridgeless) interpolant—matching the **double-descent** pattern we observed.
- Natural question: **Can we approximate deep nets as linear or kernel machines** so we can **port these insights** (interpolation, minimum-norm bias, double descent) to modern networks?
- Idea: look at **very wide networks trained with small steps**; their training can be **linearized around initialization** and viewed through a **kernel** induced by the network itself.
- This leads us to the **Neural Tangent Kernel (NTK)**—a framework where wide nets behave like **kernel regressors** with a learned, network-defined kernel.
- **Next:** define NTK, show the training dynamics, and connect it back to interpolation and generalization.

# Linearization of Neural Networks

***For very wide nets trained with small steps, training behaves like fitting a linear model on fixed features determined at initialization.***

Setup (scalar output for simplicity)

- **Data:**  $X = \{x_i\}_{i=1}^n$ , with  $x_i \in \mathbb{R}^d$ ; labels  $y \in \mathbb{R}^n$  where  $y_i \in \mathbb{R}$ .
- **Network:**  $f_\theta : \mathbb{R}^d \rightarrow \mathbb{R}$  with parameters  $\theta \in \mathbb{R}^p$ .
- **Initialization:**  $\theta_0 \in \mathbb{R}^p$ ; **parameter change:**  $\Delta\theta := \theta - \theta_0$ .

Fixed "tangent" features at initialization

- $\phi(x)$ : gradient of the output w.r.t. parameters, evaluated at  $\theta_0$  (a length- $p$  feature vector).

$$\phi(x) := \nabla_\theta f_{\theta_0}(x) \in \mathbb{R}^p.$$

First-order (linear) approximation around  $\theta_0$

$$f_\theta(x) \approx f_{\theta_0}(x) + \phi(x)^\top \Delta\theta$$

- $f_{\theta_0}(x)$ : prediction at initialization.
- $\phi(x)^\top \Delta\theta$ : linear model in the **fixed features**  $\phi(x)$ .

***Wide + small steps  $\Rightarrow$  NN  $\approx$  NTK kernel regression  $\Rightarrow$  double descent/benign overfitting results from linear & kernel models largely apply.***

Neural Tangent Kernel (NTK)

$$K(x, x') := \phi(x)^\top \phi(x') \in \mathbb{R}.$$

- **Kernel Gram matrix (train–train):**  $K(X, X) \in \mathbb{R}^{n \times n}$  with entries  $K_{ij} = K(x_i, x_j)$ .
- **Cross-kernel (test–train):**  $K(x, X) \in \mathbb{R}^n$  with entries  $K(x, x_i)$ .

Training dynamics (squared loss, gradient flow)

Let  $f_t(X) \in \mathbb{R}^n$  collect predictions  $(f_{\theta(t)}(x_1), \dots, f_{\theta(t)}(x_n))$ .

Under gradient flow (continuous-time limit of small-step GD) in the lazy regime:

$$\frac{df_t(X)}{dt} = -K(X, X) (f_t(X) - y)$$

- $t$ : (continuous) training time.
- Interpretation: predictions evolve by **kernel gradient descent** with kernel  $K$ .

Closed-form limit (interpolation,  $K(X, X)$  invertible)

$$f_\infty(x) = K(x, X) K(X, X)^{-1} y$$

- The trained wide NN behaves like **kernel regression with the NTK**.
- If  $K(X, X)$  is ill-conditioned, use ridge:

$$f_\lambda(x) = K(x, X) (K(X, X) + \lambda I)^{-1} y, \quad \lambda > 0.$$

# Implicit Bias of Gradient Descent

- **On linearly separable data, plain gradient descent on logistic/cross-entropy (and any smooth, strictly decreasing loss with an exponential tail) implicitly drives the predictor toward the hard-margin SVM solution.** Concretely, the weights blow up in norm but their **direction** converges to the L2 max-margin separator.
- Under similar assumptions, **SGD also converges in direction to the L2 max-margin separator** on linearly separable data with exponentially-tailed losses (e.g., logistic).
- **GD/SGD pick a specific interpolating classifier:** on linearly separable data with an exponentially-tailed loss (e.g., logistic), both **gradient descent** and **SGD** drive the weight norm to infinity but their **direction** converges to the **L2 hard-margin SVM**. Practically, training longer after achieving 0 training error keeps **increasing the margin**, which often improves test accuracy.
  - If you're doing linear (or last-layer) classification with cross-entropy, **plain GD/SGD implicitly does SVM-like margin maximization**—so longer training can help margins without explicit regularization.

Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, Nathan Srebro (2018), [The implicit Bias of Gradient Descent on Separable Data](#)

Mor Shpigel Nacson et al (2019), [Stochastic Gradient Descent on Separable Data: Exact Convergence with a Fixed Learning Rate](#)

# Taxonomy of overfitting




Figure 1: As  $n \rightarrow \infty$ , interpolating methods can exhibit three types of overfitting. (A) In *benign overfitting*, the predictor asymptotically approaches the ground-truth, Bayes-optimal function. Nadaraya-Watson kernel smoothing with a singular kernel, shown here, is asymptotically benign. (B) In *tempered overfitting*, the regime studied in this work, the predictor approaches a constant test risk greater than the Bayes-optimal risk. Piecewise-linear interpolation is asymptotically tempered. (C) In *catastrophic overfitting*, the predictor generalizes arbitrarily poorly. Rank- $n$  polynomial interpolation is asymptotically catastrophic.

Benign	Tempered	Catastrophic
<ul style="list-style-type: none"> <li>• Early-stopped DNNs</li> <li>• KR with ridge</li> <li>• <math>k</math>-NN (<math>k \sim \log n</math>)</li> <li>• Nadaraya-Watson kernel smoothing with Hilbert kernel</li> </ul>	<ul style="list-style-type: none"> <li>• <b>Interpolating DNNs</b></li> <li>• <b>Laplace KR</b></li> <li>• <b>ReLU NTKs</b></li> <li>• <math>k</math>-NN (constant <math>k</math>)</li> <li>• Simplicial interpolation</li> </ul>	<ul style="list-style-type: none"> <li>• <b>Gaussian KR</b></li> <li>• Critically-parameterized regression</li> </ul>

Table 2: A taxonomy of models under the three types of fitting identified in this work. **BOLD** are results from our work, others are known or folklore results.

	Regression	Classification
<b>Benign</b>	$\lim_{n \rightarrow \infty} \mathcal{R}_n = R^*$	$\lim_{n \rightarrow \infty} \mathcal{R}_n = R^*$
<b>Tempered</b>	$\lim_{n \rightarrow \infty} \mathcal{R}_n \in (R^*, \infty)$	$\lim_{n \rightarrow \infty} \mathcal{R}_n \in (R^*, 1 - \frac{1}{K})$
<b>Catastrophic</b>	$\lim_{n \rightarrow \infty} \mathcal{R}_n = \infty$	$\lim_{n \rightarrow \infty} \mathcal{R}_n = 1 - \frac{1}{K}$

# Deep Neural Network Diagnostics via Random Matrix Theory

## SETOL + WeightWatcher — one-slide takeaway

**Main message.** Deep nets that *generalize* have **heavy-tailed** layer spectra. You can **score a model's quality without any data** by fitting a power law to each layer's empirical spectral density (ESD) and reading off the  **$\alpha$  (alpha)** exponent; SETOL provides the theory, WeightWatcher gives you the button.

arXiv +1

**Key result (paper).** SETOL derives a **layer-quality metric** via a one-step **Exact Renormalization Group (ERG)** construction and shows it **aligns closely with HTSR's  $\alpha/\hat{\alpha}$  metrics** computed from the ESDs of weight matrices—yielding **conditions for “ideal learning”** and flags for non-ideal regimes (e.g., **overfit layers when  $\alpha < 2$** ). arXiv

**What to remember.**

- **Good layers:**  $\alpha \approx 2$  (practically 2–6 acceptable).
- **Red flags:**  $\alpha < 2$  (over-trained/memorizing) or  $\gg 6$  (under-trained/noisy); watch for **correlation traps** (spikes/outliers). WeightWatcher +1
- **Why it matters:** Average  $\hat{\alpha}$  (and related summaries) **track test accuracy trends across model families**—no train/test data needed. Nature



# References

- Mikhail Belkin, Daniel Hsu, Siyuan Ma, Soumik Mandal, [Reconciling modern machine-learning practice and the classical bias-variance trade-off](#)
- Arthur Jacot et al (2020), [Neural Tangent Kernel: Convergence and Generalization in Neural Networks](#)
- Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, Nathan Srebro (2018), [The implicit Bias of Gradient Descent on Separable Data](#)
- Mor Shpigel Nacson et al (2019), [Stochastic Gradient Descent on Separable Data: Exact Convergence with a Fixed Learning Rate](#)
- Neil Mallinar et al (2024), [Benign, Temered, or Catastrophic: A taxonomy of Overfitting](#)
- Charles Martin and Michael Mahoney (2021), [Implicit self-regularization in deep neural networks](#)
- Oualid Missaoui (2025), [Inductive Triplet Fine Tuning for Small Language Models](#)